



BY Developers FOR Developers

Storage Developer Conference
September 22-23, 2020

pynvme: an open, fast and extensible NVMe SSD test tool

Crane Chu, Engineer, Founder
GENG YUN Technology Pte Ltd





Requirement

Changes in SSD

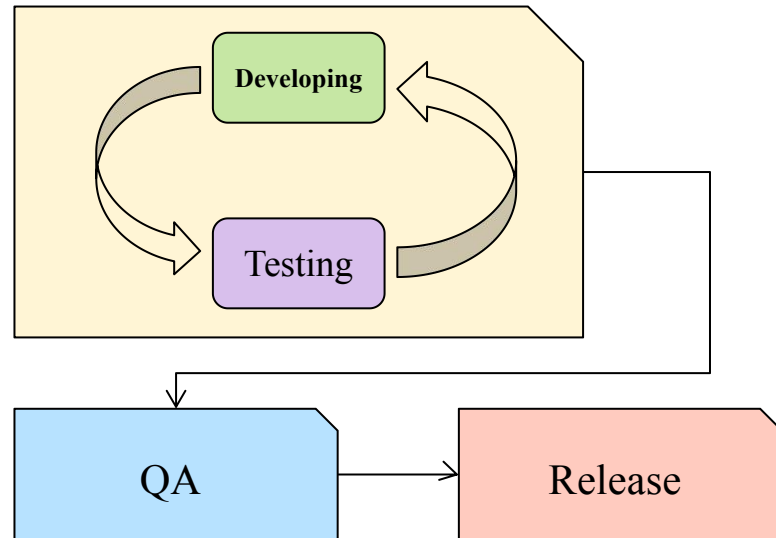
- SSD has been changing for the decade:
 - media:
 - SLC, MLC, TLC, QLC
 - 2D => 3D
 - PCM, 3D-Xpoint...
 - host:
 - PATA, SATA, PCIe/NVMe
 - open-channel
 - up coming: ZNS, KV, ... ?
 - DRAM
 - form factor
- Agile: good for the constant change and uncertainty



Agile Testing

- Developing and Testing are done interactively and iteratively.
- QA verifies the product of Dev/Test for customers.
- Testing and QA are different. Testing tools and QA tools are also different.
- Most available tools in the market are QA tools.

test	QA
for developer	for customer
before checkin	before release
automatic	manual
white-box	black-box
short	long
changing	stable



Experience with dnvme

- extended dnvme with a python wrapper in user space, thus developers can write test scripts in Python.
- Some essential problems:
 - performance:
 - IOPS, latency, consistency
 - test efficiency
 - stress tests
 - maintainness: kernel module
 - memory: user can only allocate virtual memory
- Then, SPDK comes ...
 - high performance
 - user space driver
 - DMA memory

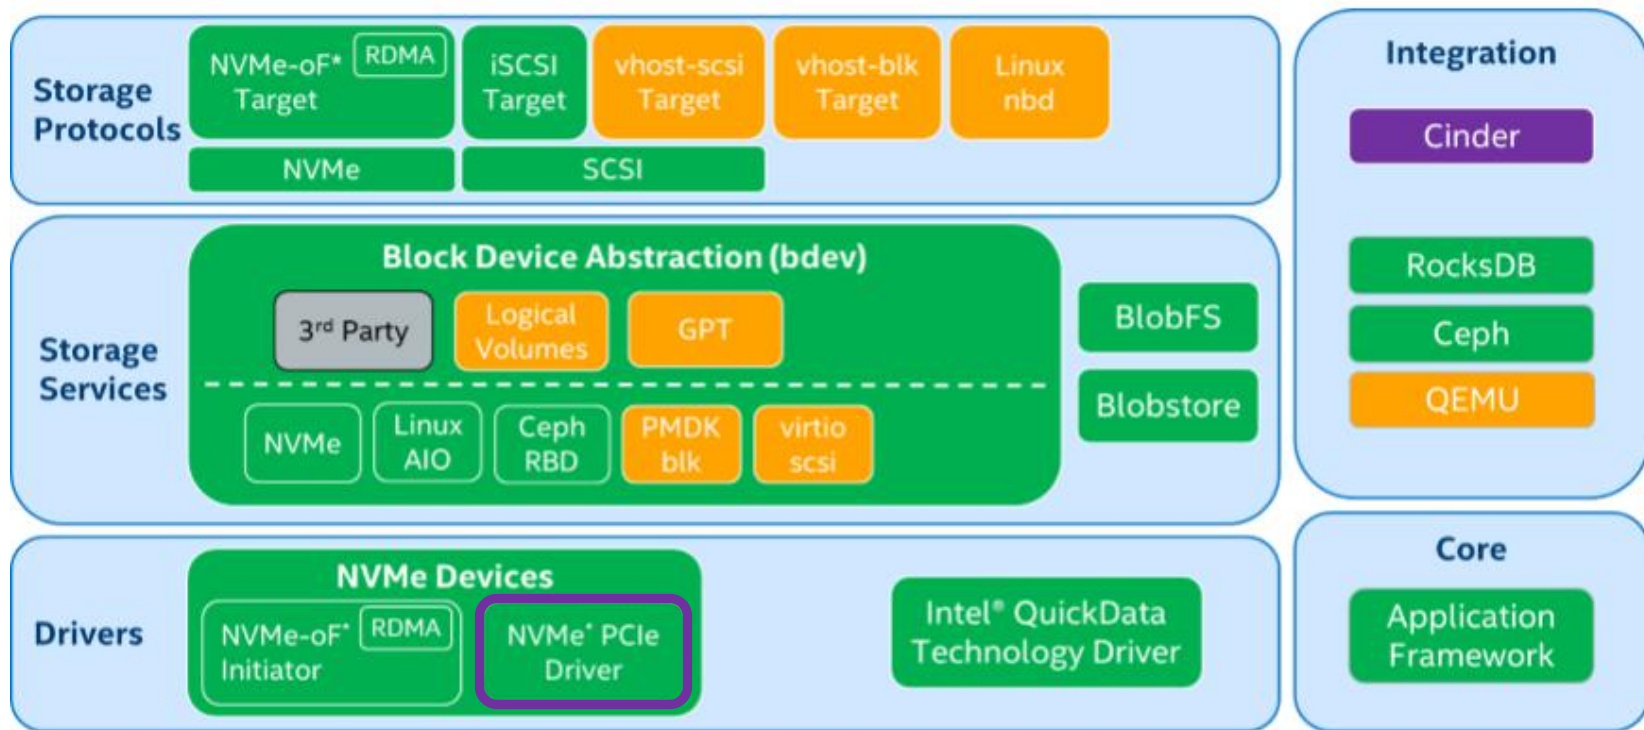


- python + nvme: pynvme <https://github.com/pynvme/pynvme>



Design

SPDK



Architecture

library (ongoing):

ZNS

TCG

psd

API:

controller

namespace

qpair

pcie

buffer

SPDK: nvme driver

cmdlog

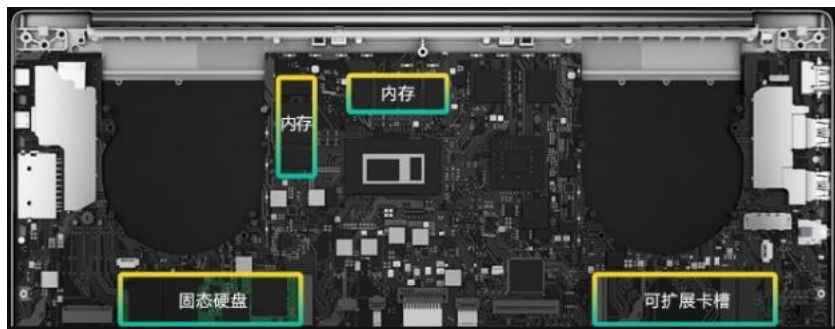
checksum

ioworker

Open Ecosystem



Hardware



High Performance

fio (unit: K IOPS)

Q count	1	2	4
test 1	200	332	353
test 2	211	319	340
test 3	211	248	354

pynvme (unit: K IOPS)

Q count	1	2	4
test 1	358	359	359
test 2	358	358	359
test 3	358	356	359



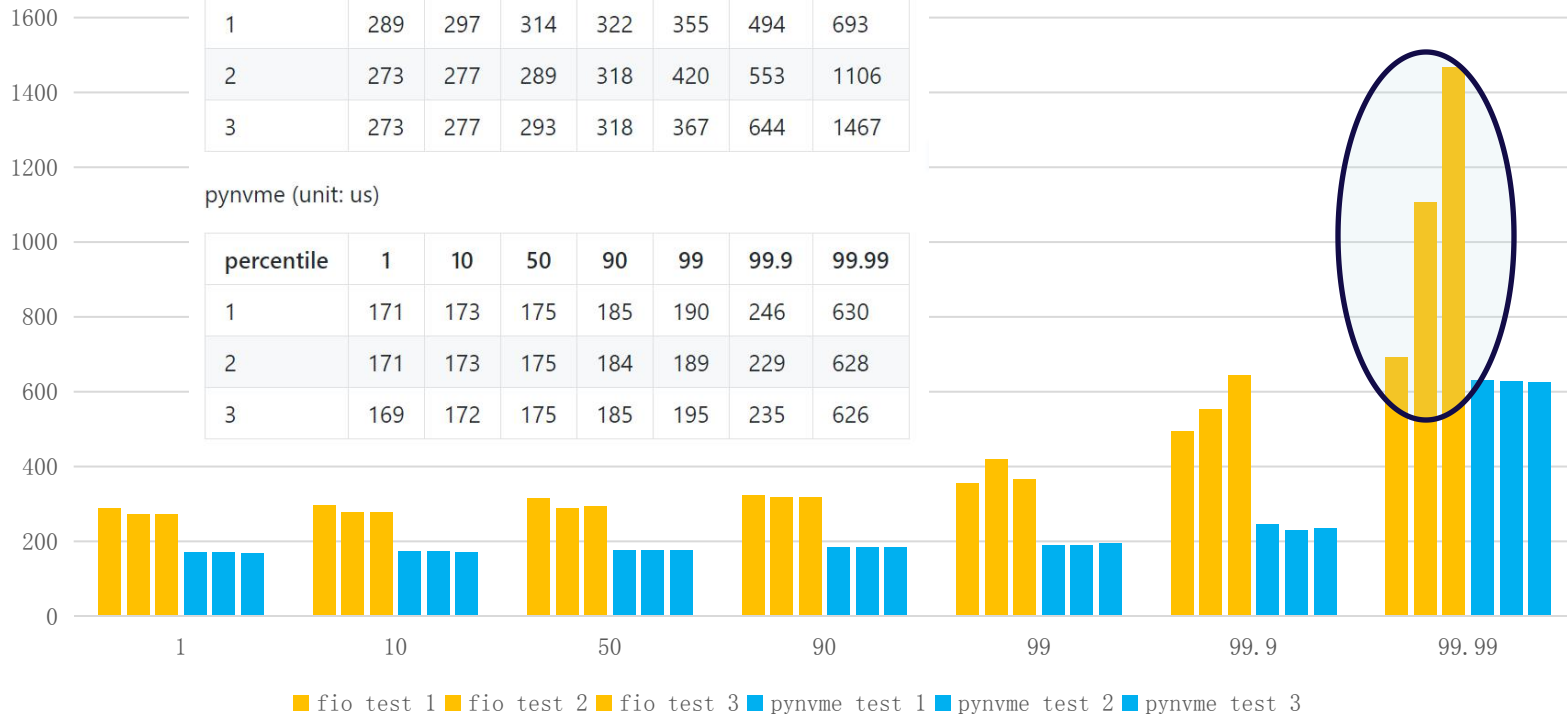
Low Latency

fio (unit: us)

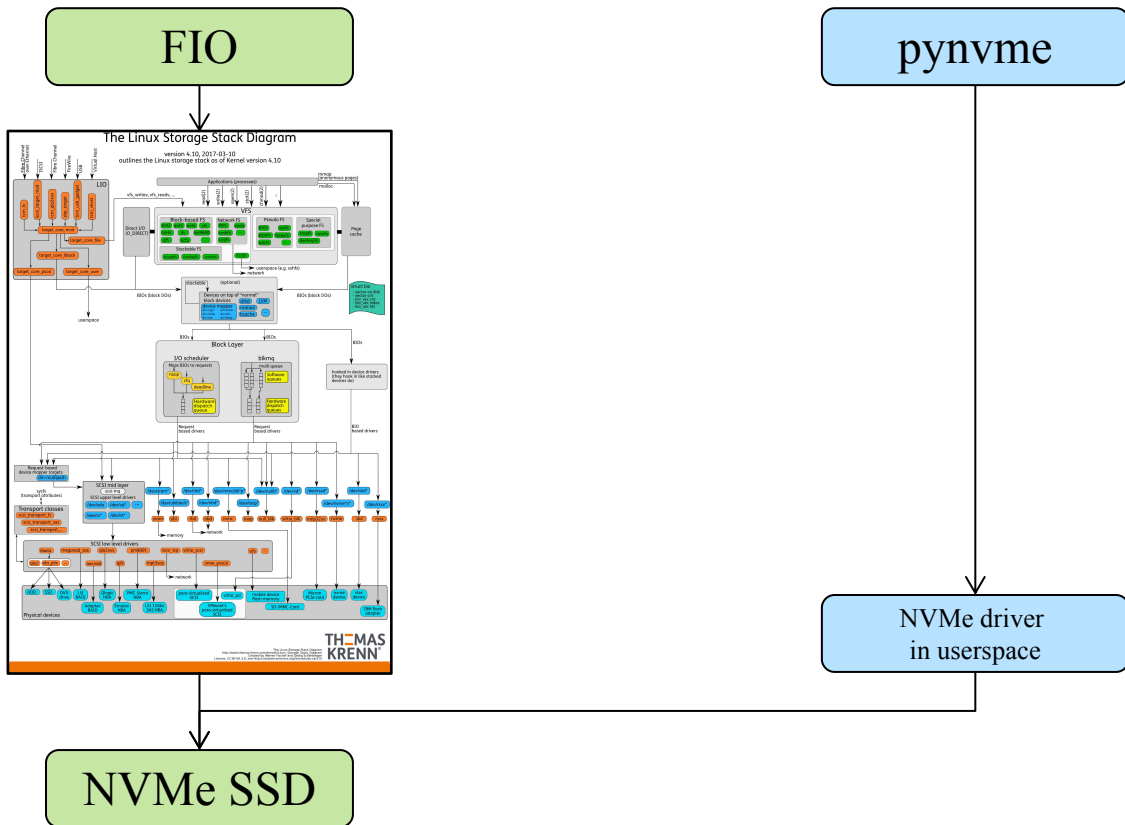
percentile	1	10	50	90	99	99.9	99.99
1	289	297	314	322	355	494	693
2	273	277	289	318	420	553	1106
3	273	277	293	318	367	644	1467

pynvme (unit: us)

percentile	1	10	50	90	99	99.9	99.99
1	171	173	175	185	190	246	630
2	171	173	175	184	189	229	628
3	169	172	175	185	195	235	626



Design



psd: Python Space Driver

psd:

IOSQ

IOCQ

SQE

CQE

PRPList

API:

controller

namespace

qpair

pcie

buffer

SPDK: nvme driver

cmdlog

checksum

ioworker

NVMe SSD

Features

- access PCI configuration space
- access NVMe registers in BAR space
- send any NVMe admin/IO commands
- support callback functions for NVMe commands
- support MSI/MSIx interrupts
- transparent checksum verification on every LBA
- generates IO workload of high performance and low latency
- support multiple namespaces
- support multiple tests on different controllers
- integrate with the test framework pytest
- integrate with VSCode to display cmdlog in GUI
- support NVMe over TCP targets
- doc: <https://pynvme.readthedocs.io/>



Examples

Example: hello world

```
1 import time
2 import pytest
3 import logging
4
5 import nvme as d
6
7
8 # intuitive, spec, qpair, vscode, debug, cmdlog, assert
9 def test_hello_world(nvme0, nvme0n1, qpair):
10     # prepare data buffer and IO queue
11     read_buf = d.Buffer(512)
12     write_buf = d.Buffer(512)
13     write_buf[10:21] = b'hello world'
14
15     # send write and read command
16     def write_cb(cdw0, status1): # command callback function
17         nvme0n1.read(qpair, read_buf, 0, 1)
18     nvme0n1.write(qpair, write_buf, 0, 1, cb=write_cb)
19
20     # wait commands complete and verify data
21     assert read_buf[10:21] != b'hello world'
22     qpair.waitdone(2)
23     assert read_buf[10:21] == b'hello world'
```

- pytest fixtures
 - callback functions for commands
 - sync point: waitdone()
-
- run test:
 - > make setup
 - > make test TESTS=scripts/test_examples.py::test_hello_world
 - > make test TESTS=scripts/test_examples.py::test_hello_world pciaddr=0000:3d:00.0
 - > make test TESTS=scripts/test_examples.py::test_hello_world pciaddr=172.168.5.44

Example: psd

```
698 def test_write_before_power_cycle(nvme0, subsystem):
699     cq = IOCQ(nvme0, 1, 128, PRP(2*1024))
700     sq = IOSQ(nvme0, 1, 128, PRP(8*1024), cqid=1)
701
702     #burst write
703     for i in range(127):
704         cmd = SQE(1, 1)
705         buf = PRP(512, ptype=32, pvalue=i)
706         cmd.prp1 = buf
707         cmd[10] = i
708         sq[i] = cmd
709
710     # write 127 512byte at one shot
711     sq.tail = 127
```

- create IOCQ/IOSQ respectively with specified sqid and cqid
- fill commands in SQ entries
- manually trigger doorbell

Example: ioworker

```
754 def test_ioworker_single(nvme0n1):
755     nvme0n1.ioworker(io_size=2, time=2).start().close()
756
757
758 def test_ioworker_multiple(nvme0n1, qpair):
759     with nvme0n1.ioworker(io_size=8, lba_random=False,
760                           read_percentage=100, time=2), \
761         nvme0n1.ioworker(io_size=64, lba_random=True,
762                           read_percentage=50, time=2):
763         pass
764
```

- fio-like IO generator
 - Python API
 - better performance
- sending IO in separated processes
- define IO pattern in parameters
- support multiple ioworkers
- support workload from simple to complex

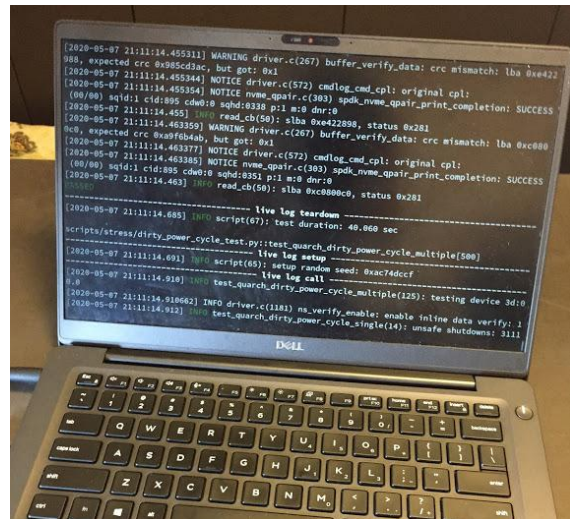
Example: dirty power cycle

```

33      # 128K random write
34      cmdlog_list = [None]*1000
35      with nvme0n1.ioworker(io_size=256,
36                             lba_random=True,
37                             read_percentage=30,
38                             region_end=region_end,
39                             time=30,
40                             qdepth=qdepth,
41                             output_cmdlog_list=cmdlog_list):
42          # sudden power loss before the ioworker end
43          time.sleep(10)
44          subsystem.poweroff()
45
46      # power on and reset controller
47      time.sleep(5)
48      subsystem.poweron()
49      nvme0.reset()

```

- ioworker sending IO in a separated process
- cut power in main process when ioworker is working



Example: customized nvme init

```
755 def test_weighted_round_robin_arbitration(pcie):
756     def nvme_init_wrr(nvme0):
757         # disable cc.en
758         nvme0[0x14] = 0
759         while not (nvme0[0x1c]&0x1) == 0: pass
760         nvme0.init_adminq()
761
762         # enable WRR and cc.en
763         nvme0[0x14] = 0x00460801
764         while not (nvme0[0x1c]&0x1) == 1: pass
765
766         nvme0.identify(d.Buffer(4096)).waitdone()
767         nvme0.init_ns()
768         nvme0.setfeatures(0x7, cdw11=0xffffefffe).waitdone()
769         cdw0 = nvme0.getfeatures(0x7).waitdone()
770         nvme0.init_queues(cdw0)
771
772     nvme0 = d.Controller(pcie, nvme_init_func=nvme_init_wrr)
773     if (nvme0.cap>>17) & 0x1 == 0:
774         pytest.skip("WRR is not supported")
775
776     # set arbitration weight
777     assert nvme0[0x14] == 0x00460801
778     nvme0.setfeatures(1, cdw11=0x07030103).waitdone()
779     cdw0 = nvme0.getfeatures(1).waitdone()
780     assert cdw0 == 0x07030103
781
```

- pynvme has a default nvme initializaiton process
- users can define different nvme initialization process in test scripts

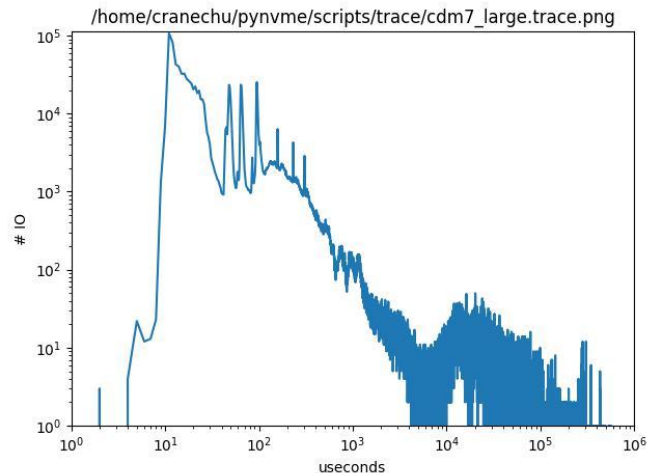
Example: latency of JEDEC workload

```

794 def test_ioworker_jedec_enterprise_workload_512(nvme0n1):
795     distribution = [1000]*5 + [200]*15 + [25]*80
796     iosz_distribution = {1: 4,
797                          2: 1,
798                          3: 1,
799                          4: 1,
800                          5: 1,
801                          6: 1,
802                          7: 1,
803                          8: 67,
804                          16: 10,
805                          32: 7,
806                          64: 3,
807                          128: 3}
808
809     output_percentile_latency = dict.fromkeys([99, 99.99, 99.9999])
810     nvme0n1.ioworker(io_size=iosz_distribution,
811                      lba_random=True,
812                      qdepth=128,
813                      distribution = distribution,
814                      read_percentage=0,
815                      ptype=0xbeef, pvalue=100,
816                      time=30,
817                      output_percentile_latency=\
818                          output_percentile_latency).start().close()
819     logging.info(output_percentile_latency)
820

```

- ioworker generates IO according to JEDEC enterprise workload
- scripts get latency at different percentile (99%, 99.99%, 99.9999%)



Example: check on error code

```
42 def test_read_large_lba(nvme0, nvme0n1, buf, qpair):
43     ncap = nvme0n1.id_data(15, 8)
44
45     nvme0n1.read(qpair, buf, ncap-1).waitdone()
46     with pytest.warns(UserWarning, match="ERROR status: 00/80"):
47         nvme0n1.read(qpair, buf, ncap).waitdone()
48     with pytest.warns(UserWarning, match="ERROR status: 00/80"):
49         nvme0n1.read(qpair, buf, ncap+1).waitdone()
50     with pytest.warns(UserWarning, match="ERROR status: 00/80"):
51         nvme0n1.read(qpair, buf, ncap-1, 2).waitdone()
52     with pytest.warns(UserWarning, match="ERROR status: 00/80"):
53         nvme0n1.read(qpair, buf, 0xffffffff00000000).waitdone()
54
```

- pynvme through a warning when command completes with an error
- scripts can use pytest to capture and check the warning

Example: sanitize an aer

```
93 def test_write_in_sanitize_operations(nvme0, nvme0n1, buf, qpair):
94     if nvme0.id_data(331, 328) == 0: #L9
95         pytest.skip("sanitize operation is not supported") #L10
96
97     logging.info("supported sanitize operation: %d" % nvme0.id_data(331, 328))
98     nvme0.sanitize().waitdone() #L13
99
100     with pytest.warns(UserWarning, match="ERROR status: 00/1d"):
101         nvme0n1.write(qpair, buf, 0).waitdone()
102
103     # check sanitize status in log page
104     with pytest.warns(UserWarning, match="AER notification is triggered"):
105         nvme0.getlogpage(0x81, buf, 20).waitdone() #L17
106         while buf.data(3, 2) & 0x7 != 1: #L18
107             time.sleep(1)
108             nvme0.getlogpage(0x81, buf, 20).waitdone() #L20
109             progress = buf.data(1, 0)*100//0xffff
110             logging.info("%d%%" % progress)
111
112     # check sanitize status
113     nvme0.getlogpage(0x81, buf, 20).waitdone()
114     assert buf.data(3, 2) & 0x7 == 1
115
```

- check if sanitize is supported
- start sanitize operation
- send a write command, and check if it is aborted due to the sanitize operation in progress
- monitor the sanitize progress till it is completed
- check the final sanitize status in logpage

Example: multiple processes

```
774 def test_ioworker_with_temperature(nvme0, nvme0n1, buf):
775     with nvme0n1.ioworker(io_size=256,
776                           time=30,
777                           op_percentage={0:10, # flush
778                                           2:60, # read
779                                           9:30}), \
780         nvme0n1.ioworker(io_size=8,
781                           time=30,
782                           op_percentage={0:10, # flush
783                                           9:10, # trim
784                                           1:80}):# write
785         for i in range(40):
786             time.sleep(1)
787             nvme0.getlogpage(0x02, buf, 512).waitdone()
788             ktemp = buf.data(2, 1)
789             from pytemperature import k2c
790             logging.info("temperature: %.2f degreeC" %
791                          k2c(ktemp))
792
```

- process 1: ioworker with flush, read, trim commands
- process 2: ioworker with flush, write, trim commands
- main process: print temperature value in SMART data for every second

Live Demo in VSCode

The screenshot displays the VS Code interface during a live demo. The left sidebar shows the file explorer with a project named 'TEST'. The code editor displays a file named 'test_examples.py' with a line of code: 'CAZ-82256-Q11 NVMe LITEON 256GB'. The terminal window shows the output of a command, including a timestamp and a list of data points. The Python Test Log window shows a series of log entries, including 'INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC'. At the bottom, two performance gauges are visible: 'pynvme' showing a value of 219.26 MB/s and 'IOPS' showing a value of 003.25 K/s.

TEST

test_examples.py

CMDLOG 0000:3d:00.0 Q01 X

1 2020-08-25 17:11:15.660273 [cmd001: Dataset Mar

2 0x00100009, 0x00000001, 0x00000000, 0x00000000

3 0000:3d:00.0 Q01: >>>> 0000, 0x3ebff000, 0x00000000

4 0000:3d:00.0 Q02: >>>> 0000, 0x00000000, 0x00000000

5 0x00000000, 0x00000000, 0x00000000, 0x00000000

6 not completed

7 ...

8

9 2020-08-25 17:11:15.660268 [cmd002: Read]

10 0x00030002, 0x00000001, 0x00000000, 0x00000000

11 0x00000000, 0x00000000, 0x3ef3f000, 0x00000000

12 0x3edc3038, 0x00000000, 0x0a9a793c, 0x00000000

13 0x000000ff, 0x00000000, 0x00000000, 0x00000000

14 not completed

15 ...

16

17 2020-08-25 17:11:15.660261 [cmd003: Dataset Mar

18 0x00160009, 0x00000001, 0x00000000, 0x00000000

19 0x00000000, 0x00000000, 0x3ea7f000, 0x00000000

20 0x00000000, 0x00000000, 0x00000000, 0x00000004

21 0x00000000, 0x00000000, 0x00000000, 0x00000000

22 not completed

23 ...

Python Test Log

2020-08-25 17:11:09.715] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:09.715] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:10.754] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:11.796] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:12.835] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:13.883] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:14.920] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:15.964] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:17.084] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:18.123] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:19.361] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:20.400] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:22.136] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:23.174] INFO test_ioworker_with_temperature (790): temperature: 55.85 degreeC

2020-08-25 17:11:24.206] INFO test_ioworker_with_temperature (790): temperature: 56.85 degreeC

2020-08-25 17:11:25.243] INFO test_ioworker_with_temperature (790): temperature: 56.85 degreeC

2020-08-25 17:11:26.298] INFO test_ioworker_with_temperature (790): temperature: 56.85 degreeC

2020-08-25 17:11:27.348] INFO test_ioworker_with_temperature (790): temperature: 56.85 degreeC

Performance Gauge X

pynvme

219.26 MB/s

IOPS

003.25 K/s

master* Python 3.8.5 64-bit 7 Running Tests

Ln 1, Col 1 Spaces: 4 Plain Text

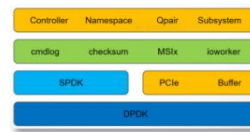
pynvme builds your own tests.



Ecosystem



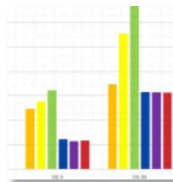
Scripts



Extendability



Hardware

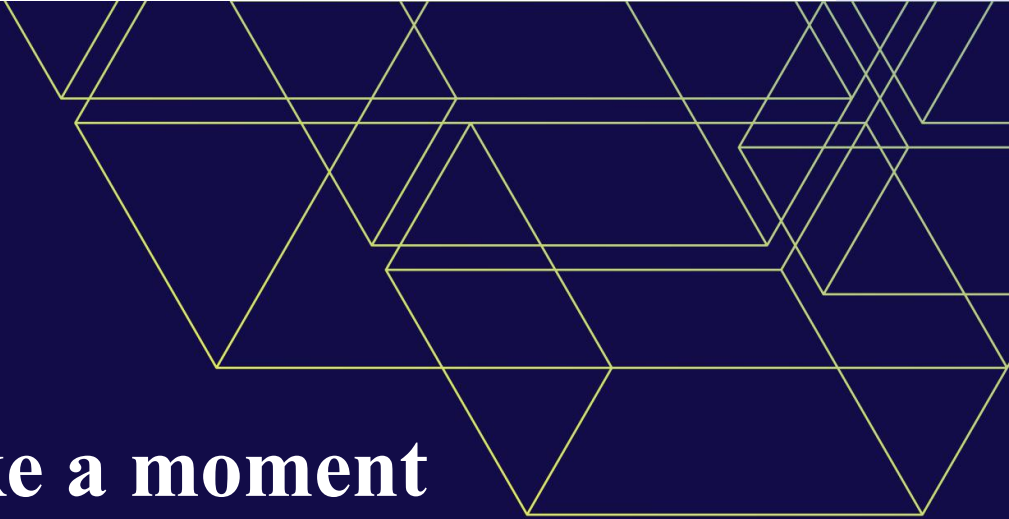


Performance



Service





**Please take a moment
to rate this session.**

Your feedback matters to us.

Thank you!!!